

# FlyClient Guide

Succinct chain verification: the sampling protocol, the deployed ZIP-221 commitment, and the unbuilt bridge

m@rek.onl

## Abstract

Assuming the deployed chain and header rules of the *Ironwood Guide* and the light-client layer of the *Sync Guide*, this volume of *The Zcash Arboretum* documents succinct chain verification for Zcash: the FlyClient protocol as its paper constructs it, the ZIP-221 chain-history commitment that every Zcash block header has carried since Heartwood, and the distance between the two. The commitment is consensus-critical and universally maintained; no deployed light client consumes it. Claims about the deployed tree cite the ZIPs and the implementations; claims about the sampling protocol carry the paper's own model assumptions, and its printed defects are registered rather than silently repaired; the bridge between the two is classified plainly as unbuilt.

# Contents

|          |                                                                 |           |
|----------|-----------------------------------------------------------------|-----------|
| <b>1</b> | <b>Scope, sources, and the deployment boundary</b>              | <b>4</b>  |
| 1.1      | Sources and pinning . . . . .                                   | 4         |
| 1.2      | Relation to the sibling volumes . . . . .                       | 4         |
| 1.3      | The deployment boundary, stated at the outset . . . . .         | 5         |
| <b>2</b> | <b>The problem: verifying a chain without downloading it</b>    | <b>5</b>  |
| 2.1      | The cost of header-only verification . . . . .                  | 5         |
| 2.2      | Variable difficulty and the difficulty-raising attack . . . . . | 5         |
| 2.3      | Prior art: superblocks and their failure modes . . . . .        | 6         |
| <b>3</b> | <b>Merkle mountain ranges</b>                                   | <b>6</b>  |
| 3.1      | The structure . . . . .                                         | 6         |
| 3.2      | Appending in logarithmic time . . . . .                         | 7         |
| 3.3      | The prefix-extension trick . . . . .                            | 8         |
| 3.4      | Headers as chain commitments . . . . .                          | 8         |
| 3.5      | The difficulty MMR . . . . .                                    | 8         |
| 3.6      | Invalid transitions never help . . . . .                        | 9         |
| <b>4</b> | <b>The FlyClient protocol</b>                                   | <b>10</b> |
| 4.1      | The model, and what the adversary may do . . . . .              | 10        |
| 4.2      | What security means here . . . . .                              | 11        |
| 4.3      | The interactive game . . . . .                                  | 11        |
| 4.4      | From uniform sampling to the optimal distribution . . . . .     | 11        |
| 4.5      | Removing the verifier . . . . .                                 | 13        |
| 4.6      | The theorem, and staying synchronised . . . . .                 | 13        |
| 4.7      | What the evaluation showed . . . . .                            | 13        |
| <b>5</b> | <b>ZIP-221: the deployed chain-history tree</b>                 | <b>14</b> |
| 5.1      | Epoch scoping: one tree per network upgrade . . . . .           | 14        |
| 5.2      | Numbering, peaks, and bagging . . . . .                         | 15        |
| 5.3      | Node metadata: the complete field list . . . . .                | 15        |
| 5.4      | Hashing and personalisation . . . . .                           | 15        |
| 5.5      | The consensus rules at Heartwood . . . . .                      | 16        |
| 5.6      | Deltas from the paper’s construction . . . . .                  | 17        |
| <b>6</b> | <b>NU5: the commitment becomes one of two</b>                   | <b>17</b> |
| 6.1      | Why an authorising-data root exists . . . . .                   | 17        |
| 6.2      | hashBlockCommitments . . . . .                                  | 18        |
| 6.3      | Activation constants . . . . .                                  | 18        |
| <b>7</b> | <b>Deployed implementations</b>                                 | <b>18</b> |
| 7.1      | The zcash_history crate . . . . .                               | 19        |
| 7.2      | Zebra . . . . .                                                 | 19        |
| 7.3      | zcashd . . . . .                                                | 20        |
| 7.4      | Divergence register . . . . .                                   | 20        |
| <b>8</b> | <b>The deployment gap</b>                                       | <b>21</b> |
| 8.1      | Who computes the root, who verifies it . . . . .                | 21        |
| 8.2      | The light-client stack consumes none of it . . . . .            | 21        |
| 8.3      | What an actual Zcash FlyClient still needs . . . . .            | 22        |
| 8.4      | Elsewhere: velvet paths and inspired bridges . . . . .          | 22        |

|     |                                    |    |
|-----|------------------------------------|----|
| 8.5 | Relation to the frontier . . . . . | 23 |
| 8.6 | Closing the boundary . . . . .     | 23 |

# 1 Scope, sources, and the deployment boundary

FlyClient is a light-client protocol that verifies a proof-of-work chain by sampling logarithmically many block headers under a difficulty-weighted distribution, checking each sample against a commitment that every header carries over the entire chain before it (Bünz–Kiffer–Luu–Zamani, “FlyClient: Super-Light Clients for Cryptocurrencies”, IACR ePrint 2019/226; published at IEEE S&P 2020). The commitment is a Merkle mountain range — an append-friendly Merkle tree — and the protocol’s claim is quantitative: a client holding only the genesis block can accept the head of an  $n$ -block chain after downloading  $O(\lambda \log_{1/c}(n) \log_2(n) + L)$  data rather than  $n$  headers, where  $c$  bounds the adversary’s mining power relative to the honest network’s and  $L$  is the fork length below which the adversary is unconstrained — so the client checks the last  $L$  headers in full.

This volume documents two objects and the distance between them. The first is the protocol of the paper: model, commitment, sampling distribution, security theorem, and the proof structure behind it (§2–§4). The second is the consensus structure Zcash actually deployed for it: ZIP-221’s chain-history tree, a difficulty-MMR variant whose root every Zcash block header has committed to since the Heartwood upgrade, extended at NU5 by ZIP-244’s authorising-data commitment (§5–§7). The distance is the volume’s third subject (§8): Zcash ships the server half of FlyClient in every block and validates it in every full node, yet no deployed wallet, SDK, or light-client server exposes or consumes it; deployed light clients instead trust their server for chain state outright (*Sync Guide*, §“Architecture and authority”).

## 1.1 Sources and pinning

Table 1 pins the artefacts this volume treats as ground truth. Two rigor conventions from the sibling volumes apply. Deployed behaviour cites crate, file, and line against the pinned commits; the classification of design-stage material follows the three-bin rule of the *Tachyon Guide* (§“Scope, sources, and method”): *specified*, *designed but unspecified*, or *open problem*. The FlyClient paper itself is treated as a primary source with registered defects: the revision of record contains internal inconsistencies (a sign typo in a difficulty bound, a vestigial pseudocode return, conflicting evaluation constants) that are catalogued in Remark 3.5 and flagged inline where they matter, never silently repaired.

| Artefact                          | Pin                              | Rigor                            |
|-----------------------------------|----------------------------------|----------------------------------|
| FlyClient paper                   | ePrint 2019/226, rev. 2020-08-20 | peer-reviewed; errata registered |
| ZIP 221, ZIP 244                  | zips @ 69610984 (2026-07-09)     | merged, deployed consensus       |
| zcash_history                     | 0.5.0 (crates.io)                | deployed (Zebra dependency)      |
| zebra                             | @ de14bef05 (2026-07-17)         | deployed validator               |
| zcashd                            | @ 16ac74376 (2025-10-03)         | historical validator (retired)   |
| lightwalletd, Zaino, librustzcash | repo greps, 2026-07-19           | negative results (§8)            |

Table 1: Primary sources, pinned. Verification date 2026-07-19. The paper’s local copy is `.wip/flyclient-drafts/eprint-2019-226.pdf`.

## 1.2 Relation to the sibling volumes

This volume assumes the *Ironwood Guide* for everything inside a block — notes, commitments, anchors, the NU5 transaction format — and cites its treatment of the note-commitment tree (§“Resting: the note in the commitment tree”) rather than re-deriving Merkle mechanics; the general theory of commitments and collision resistance is the *Crypto Guide*’s (§“Commitment schemes”). The *Sync Guide* owns the deployed light-client layer whose central trust assumption — the node/proxy combination “is trusted to provide a correct view of the current best chain state” (ZIP 307, §“The stated security model”) — is precisely the assumption a FlyClient bridge would discharge; this volume owns that cross-reference. Toward the frontier volumes the relationship is converse: the *Tachyon Guide*

records a draft amendment to ZIP-221 among its peripheral ZIPs, and the *Crosslink Guide*’s finality layer would change what “the best chain” means for any weight-sampling client; both interactions are taken up in §8, and per the sibling-ownership convention the later volume owns the analysis.

### 1.3 The deployment boundary, stated at the outset

Every claim in this volume sits on one side of a line that the material itself draws sharply.

- *Deployed and consensus-critical.* The ZIP-221 chain-history tree: maintained by every validator, its root committed in every block header since Heartwood (mainnet height 903,000) and folded into `hashBlockCommitments` since NU5 (mainnet height 1,687,104), validated on every block connect by zebra and, until its retirement, `zcashd` (§5–§7).
- *Specified but consumed by no one.* The same tree as a light-client instrument: ZIP-221’s stated motivation is FlyClient, and the metadata its nodes carry exists for difficulty-weighted sampling; no deployed RPC serves an MMR proof and no deployed client requests one (§8).
- *Designed but not specified for Zcash.* The FlyClient protocol itself — sampling distribution, non-interactive transcript, proof format — exists in the paper with a security theorem in the paper’s own backbone model; no ZIP specifies a Zcash instantiation, and the paper’s model does not automatically cover Zcash’s difficulty-adjustment rule (§4).

## 2 The problem: verifying a chain without downloading it

### 2.1 The cost of header-only verification

A full node verifies a proof-of-work chain by downloading and checking every block — for Bitcoin or Ethereum in 2019, more than 200 GB (ePrint 2019/226, §1). Nakamoto’s simplified payment verification reduces this to headers alone: 80 bytes rather than roughly a megabyte per Bitcoin block. The reduction is linear, not asymptotic, and the paper’s motivating arithmetic is Ethereum’s: a 508-byte header every ~15 seconds had produced, by July 2019, more than 8.2 million blocks and hence more than 3.9 GB of headers that an SPV client must fetch and store before verifying its first payment (§1, p. 3). A device class that cannot bear this — phones, wearables, embedded clients — is exactly the class light clients exist for, and the practical alternative, a trusted wallet provider, “negates much of the benefits of these decentralized ledgers” (§1, p. 2).

An SPV client’s trust model is already weaker than a full node’s. It cannot check transaction validity or consensus rules; it assumes the most-work chain follows the rules — the paper’s Assumption 1: “The chain with the most PoW solutions follows the rules of the network and will eventually be accepted by the majority of miners” (§1, p. 2). The most-work rule itself, and the probabilistic guarantee it purchases, are the *Consensus Guide*’s subject (§“The block tree and the most-work rule”). FlyClient keeps exactly this assumption class (strengthened quantitatively, §4.1) and attacks only the cost: can a client accept the head of the chain after sublinear communication? The obstacle is that a client that no longer checks every proof of work “can be tricked into accepting an invalid chain by a malicious prover who can precompute a sufficiently-long chain using its limited computational power” (§1, p. 3).

### 2.2 Variable difficulty and the difficulty-raising attack

Any sampling client must contend with an attack that does not exist at fixed difficulty. Under variable difficulty the chain’s total *work*, not its length, decides validity, and an adversary may mine few but heavy blocks: Bahack’s difficulty-raising attack. The paper’s worked figure makes the threat concrete: an adversary holding one third of the honest mining power, permitted a single block of arbitrarily high difficulty, exceeds the expected weight of the honest chain with probability roughly 28% — nothing

close to negligible (§3.1, p. 7). Full nodes are protected by the retargeting rule itself: difficulty moves only once per epoch and only within a clamp (the dampening filter  $\tau$ ; Bitcoin operates with  $\tau = 4$ , epoch length  $m = 2016$ ). A client that samples, however, never sees most transitions, so the protocol must force *every* transition it could ever rely on to be valid — the purpose of the difficulty MMR’s metadata (§3.5) and of Lemma 4 (§3.6).

### 2.3 Prior art: superblocks and their failure modes

The prior approach to sublinear chain proofs is the superblock: a block whose hash beats a harder target than required. Superblock counting was folklore by 2012 (the “high-value-hash highway”), was made rigorous as interactive PoPoW by Kiayias–Lamprou–Stouka, and non-interactive as NIPoPoW by Kiayias–Miller–Zindros, who also broke the interactive predecessor (ePrint 2019/226, §1, p. 3). The paper’s case against the superblock family is fourfold (§1, “Current Challenges”):

1. Fixed difficulty is assumed outright, and “it isn’t clear how to modify the super-block based protocols to handle the variable difficulties.”
2. Under variable difficulty the difficulty-raising attack of §2.2 applies, and adding transition checks to superblock proofs “is a non-obvious extension.”
3. Superblocks are identifiable and carry no extra reward, so bribing and selfish-mining strategies suppress them cheaply; NIPoPoW’s published defence reverts to full SPV, so its proofs “are therefore only succinct if no adversarial influence exists.” FlyClient distinguishes blocks only by position, leaving nothing cheap to suppress.
4. NIPoPoW transaction-inclusion proofs cost roughly 15 KB on Ethereum against FlyClient’s  $\sim 1.5$  KB.

The comparison FlyClient claims for itself is summarised by the paper’s Table 1, reproduced as Table 2: three orders of magnitude against SPV at Ethereum’s 2019 length, under an adversary bounded by half the honest hash power and failure probability below  $2^{-50}$ .

| Block height   | 10 K  | 100 K  | 1,000 K | 7,000 K   |
|----------------|-------|--------|---------|-----------|
| SPV (KB)       | 4,961 | 49,609 | 496,094 | 3,472,656 |
| FlyClient (KB) | 161   | 277    | 416     | 524       |
| Improvement    | 31×   | 179×   | 1,275×  | 6,627×    |

Table 2: Proof sizes for Ethereum parameters (508-byte headers), assuming adversarial hash power at most  $c = 1/2$  of the honest hash power and failure probability below  $2^{-50}$ . Reproduced from ePrint 2019/226, Table 1 (p. 4). The 1,275× entry is the paper’s own slip: its SPV and FlyClient sizes give  $496,094/416 \approx 1,193\times$  (Remark 3.5).

One reading habit matters for every number above and below: the paper’s parameter  $c$  is the ratio of adversarial to *honest* mining power, not the adversary’s share of the total. An adversary at  $c = 0.9$  controls  $c/(1 + c) \approx 47.3\%$  of the network; the headline  $c = 1/2$  rows of Table 2 describe an adversary with a third of total power (§7.1, p. 20).

## 3 Merkle mountain ranges

### 3.1 The structure

The commitment at the heart of FlyClient is a Merkle tree variant optimised for one operation: appending a leaf without rebuilding the tree. The paper defines it recursively (ePrint 2019/226, Definition 11, p. 28); the volume restates the definition in its own numbering because everything later — ZIP-221 included — builds on it.

**Definition 3.1** (Merkle mountain range). A *Merkle mountain range* (MMR) over  $n$  leaves is a binary hash tree  $M$  with root  $r$  and depth  $\lceil \log_2 n \rceil$  such that, if  $n > 1$ , writing  $n = 2^i + j$  with  $i = \lfloor \log_2(n-1) \rfloor$  (so  $1 \leq j \leq 2^i$ ): the left child subtree of  $r$  is an MMR with  $2^i$  leaves, and the right child subtree of  $r$  is an MMR with  $j$  leaves. Each internal node holds the hash  $H(\text{left} \parallel \text{right})$  of its children’s values, for a collision-resistant  $H$ .

The left subtree is therefore the largest *perfect* binary tree the leaf count admits, and the recursion repeats on the remainder: an MMR over  $n = 6$  leaves splits  $6 = 4 + 2$  into a perfect four-leaf subtree and a perfect two-leaf subtree; an MMR over  $n = 7$  splits  $7 = 4 + 3$  and the right side splits again as  $3 = 2 + 1$ . Every MMR is a balanced binary tree in the paper’s sense (depth at most  $\lceil \log_2 n \rceil$ ), so ordinary Merkle inclusion proofs of at most  $\log_2(n) + 1$  hashes exist for every leaf (Definitions 9–10, p. 27).

**Remark 3.2** (Peaks and bagging, absent by construction). The folklore MMR presentation — Peter Todd’s, which the paper credits for the structure — keeps a *forest* of perfect trees (“peaks”), one per set bit of  $n$ , and combines (“bags”) them into a single root only when one is needed. Neither word appears anywhere in the paper: its Definition 11 is single-rooted from the start. The two presentations coincide exactly — the peaks of the classical forest are the left children of the nodes along the paper tree’s right spine together with the terminal right-spine node, and the paper’s root is the result of bagging the peaks right-to-left — but the difference in framing matters when reading ZIP-221, which specifies the same object in yet another way (§5), and when reading `zcash_history`, which stores the peaks explicitly (§7). Costs and proofs are identical between the folklore forest and the paper’s single-rooted form. ZIP-221 instead bags left-fold-leftmost (§5.2), the opposite of the paper’s implicit right fold, so with three or more peaks its bagged root and the hashes at the top of a root-anchored path differ from the paper’s; the stored mountain array and the asymptotic costs are the same.




Figure 1: The MMR over six leaves  $x_1, \dots, x_6$ , drawn in the paper’s single-rooted form (Definition 3.1):  $6 = 4 + 2$ , so the root  $r_6$  has the perfect four-leaf tree rooted at  $P_1$  as left child and the perfect two-leaf tree rooted at  $P_2$  as right child. The nodes  $P_1$  and  $P_2$  (blue fill) are exactly the *peaks* of the classical forest presentation — here the root’s left child  $P_1$  and the terminal right-spine node  $P_2$  (red edges) — and appending leaf  $x_7$  rewrites only the spine:  $P_1$  and  $P_2$  and every node beneath them are shared unchanged with  $r_7$ .

### 3.2 Appending in logarithmic time

Appending leaf  $x$  to an  $n$ -leaf MMR (the paper’s Algorithm 5, `AppendLeaf`, p. 29) has two cases. If the tree is perfect ( $n = 2^i$ ), a new root is minted with the old tree as left child and  $x$  as right child, hashing  $H(r \parallel x)$ . Otherwise the append recurses into the right subtree and rehashes on the way out. Only the right spine is ever rewritten, so the cost is  $O(\log n)$  and — decisive for a blockchain host — every perfect left subtree of the old MMR is shared, untouched, with the new one. Theorem 6 of the paper proves by induction that the result is the MMR over the old leaves with  $x$  appended rightmost.

### 3.3 The prefix-extension trick

An ordinary Merkle proof shows that a leaf lies under a root. FlyClient needs strictly more: when the verifier samples block  $k$  from a chain whose head commits to  $n$  blocks, it must check not only that block  $k$  is leaf  $k$  under the head’s root but that the MMR root stored *inside* block  $k$ ’s own header commits to exactly the first  $k - 1$  of those same leaves — that the sampled block’s view of history is a prefix of the head’s. The MMR delivers this without any additional proof data.

**Theorem 7** (ePrint 2019/226, p. 30). *For  $k \leq n$ , given  $\Pi_{x_k \in M_n}$ , the Merkle proof that leaf  $x_k$  is in the  $n$ -leaf MMR  $M_n$ , a verifier can regenerate  $r_k$ , the root of the  $k$ -leaf MMR  $M_k$ .*

The proof is an induction on the shape of Definition 3.1, and its content fits in one sentence: in this tree shape, the *left* siblings along the inclusion path of leaf  $k$  are exactly the roots of the perfect subtrees covering leaves  $1, \dots, k - 1$ , so folding them together bottom-up — the paper’s Get\_Root routine (Algorithm 6) — reconstructs the prefix root from the very path that proved inclusion. One path, two facts:  $x_k$  is the  $k$ -th leaf under  $r_n$ , and the root over the first  $k - 1$  leaves is a value the verifier can now compare against the root stored in the sampled header. The two corollaries the protocol leans on follow directly (p. 31): the MMR is position-binding for prefixes (Corollary 3: the path commits  $x_1, \dots, x_k$  as the blocks of a chain that is a prefix of the head’s), and any change to any block changes every later root (Corollary 4) — so an adversary that forks the chain can reuse no honest block from after its fork point, since the honest MMR roots inside those blocks would betray the substitution.

**Remark 3.3** (Printed defects in Algorithm 6). As printed, Get\_Root assigns from an undefined proof index ( $r = \Pi[i]$  where the loop variable is  $y$ ), ends by comparing  $y = r$  and returning a bit although its specification, its call site in Algorithm 2, and Theorem 7 all require it to return the reconstructed root, and is called with its arguments in a different order than it declares. The intended behaviour is unambiguous from Theorem 7’s proof — initialise with the first left sibling encountered, fold  $r \leftarrow H(y \parallel r)$  over the remaining left siblings bottom-up, return  $r$  — and every implementation below behaves so. Registered in Remark 3.5.

### 3.4 Headers as chain commitments

FlyClient’s consensus change is one field. At every height  $i$  the block producer appends the hash of  $B_{i-1}$  to the running MMR and records the new root  $M_i$  in the header of  $B_i$  (§2, pp. 5–6): the header of a block commits, through one MMR root, to exactly the sequence of blocks strictly before it. A verifier holding one trusted header therefore holds a commitment to the whole chain, and the per-sample check set is (Algorithm 2 with §2): the inclusion path hashes to the head’s root; the Get\_Root fold of the same path equals the root stored in the sampled header; the sampled header’s proof of work is valid.

**Remark 3.4** (Indexing drift, and the invariant that survives it). The paper is not consistent about names: §2 calls the root recorded in  $B_i$  “ $M_i$ ”, Definition 5 and Lemma 4 call the root inside the head  $B_n$  “ $M_{n-1}$ ”, and Algorithms 1–2 speak of a chain of  $n + 1$  blocks whose head commits “the first  $n$  blocks”. Every variant expresses the same invariant — *a block’s header commits to all blocks strictly before it* — and this volume uses that sentence, not any one of the indexings. ZIP-221 will shift the convention once more (the root in block  $i$  covers leaves through block  $i - 1$ , but the tree begins at an upgrade activation rather than at genesis; §5).

### 3.5 The difficulty MMR

Sampling by position is the fixed-difficulty story. Under variable difficulty the verifier must sample by *weight* — the distribution of §4.4 works over relative aggregate difficulty — and must be able to reject invalid difficulty transitions it will never individually inspect (§2.2). Both needs are served by enriching every MMR node with aggregate metadata.



**Definition 7** (Difficulty MMR; ePrint 2019/226, p. 18). *A difficulty MMR is a variant of the MMR with identical properties, but such that each node contains additional data. Every node can be written as  $\{h, w, t, D_{\text{start}}, D_{\text{next}}, n, \text{data}\}$ , where  $h$  is a  $\lambda$ -bit hash output,  $w$  is an integer total difficulty,  $t$  is a timestamp,  $D$  is a difficulty target,  $n$  is the size of the subtree rooted at the node, and  $\text{data}$  is an arbitrary string. In a leaf,  $n = 1$  and  $t$  is the time it took to mine the block;  $w, D_{\text{start}}$  are the block’s difficulty target and  $D_{\text{next}}$  the next block’s target under the difficulty-adjustment function. Each non-leaf node is  $\{H(lc, rc), lc.w + rc.w, lc.t + rc.t, lc.D_{\text{start}}, rc.D_{\text{next}}, lc.n + rc.n, \perp\}$  for children  $lc, rc$ .*

Aggregation is thus fixed once per field: weights and times sum, targets propagate from the boundary children, counts add. The verifier recomputes every parent along a sampled path from its children and checks, per level: that  $lc.D_{\text{next}} = rc.D_{\text{start}}$  (adjacent subtrees agree on the target across their boundary); that nodes below the epoch level  $\log_2 m$  carry exactly their epoch’s difficulty; that nodes spanning a transition computed  $D_{\text{next}}$  by the adjustment rule; and that a node’s aggregate  $(w, t)$  is *feasible* — some legal sequence of clamped transitions between its boundary targets could have produced them. The paper bounds feasibility explicitly: maximal weight comes from raising the difficulty by the dampening factor  $\tau$  as fast as allowed and then descending, minimal weight from the reverse; a maximal difficulty decrease needs an epoch span of at least  $\tau m/f$  rounds and a maximal difficulty increase at most  $m/(f\tau)$ . For the simplified scenario  $\tau^k D_{\text{start}} = D_{\text{next}}$  over  $n$  epochs (the paper’s own worked case, pp. 18–19) the printed upper bound is

$$w \leq D_{\text{start}} \cdot \frac{(\tau + 1) \tau^{(k+n)/2} - \tau^{1+k} - 1}{\tau - 1},$$

and the printed lower bound repeats the exponent  $(k + n)/2$  where evaluating its own summation gives  $\tau^{-(n-k)/2}$  — a sign typo, verified against the PDF at high resolution, under which the printed bound would be negative; the summation form is authoritative (Remark 3.5).

### 3.6 Invalid transitions never help

The transition checks would be pointless if an adversary could profit from difficulty rules the verifier fails to sample. The paper closes this with a rewriting argument.

**Lemma 4** (ePrint 2019/226, p. 19). *Let  $\mathcal{A}$  be an adversary in the variable-difficulty backbone model that produces, with non-negligible probability  $p$ , a chain  $C$  in which  $k$  blocks are valid. Then, assuming a collision-resistant hash function, there exists an adversary  $\mathcal{A}'$  using the same number of oracle queries that respects the retargeting rules and produces, with probability at least  $p - \text{negl}(\lambda)$ , a chain  $C'$  containing the same valid blocks.*

The proof’s mechanism is worth keeping, because it explains why the per-node checks of §3.5 suffice. Consider an invalid adjustment between  $B_i$  and  $B_{i+1}$  in  $\mathcal{A}$ ’s chain. No valid inclusion proof for  $B_i$  exists, and along any claimed path there is a highest node  $x'$  at which the verifier’s recomputation fails; every proof through  $x'$  is rejected, so *every leaf under  $x'$ ’s parent* is already unsampleable-without-rejection. The rewritten adversary  $\mathcal{A}'$  replaces the difficulty transitions inside that whole subtree with valid ones consistent with the parent’s aggregate fields — possible precisely because the parent’s own  $(w, t, D_{\text{start}}, D_{\text{next}}, n)$  pass the feasibility checks — and the blocks inside need no valid proof of work at all. Invalid transitions therefore buy nothing: the adversary might as well have used valid transitions and mined more invalid blocks, which is exactly the adversary the sampling analysis already counts.

**Remark 3.5** (Errata register for the revision of record). The ePrint revision of 2020-08-20 — the version this volume pins — contains the following printed defects, all verified against the PDF and none affecting the results: (i) the sign typo in the difficulty-MMR lower bound (§3.5); (ii) the `Get_Root` pseudocode’s undefined index, vestigial boolean return, and call-signature mismatch (Remark 3.3); (iii) two near-duplicate evaluation paragraphs in §7 carrying conflicting MMR-node overheads (8 versus 16 bytes) — the variable-difficulty evaluation is the 16-byte one; (iv) a difficulty-bomb sentence in §7.2 asserting a proof-size increase where its own mechanism and the adjacent sentence describe

a decrease; (v) AppendLeaf’s base case hashing  $H(x \parallel r)$  in prose against  $H(r \parallel x)$  in the algorithm — the algorithm is operative; (vi) Definition 7’s feasibility checks referencing  $t_{\text{start}}, t_{\text{end}}$  fields the definition never declares, and drifting between  $D_{\text{next}}$  and  $D_{\text{end}}$ ; (vii) Definition 1’s target-recalculation formula uses a symbol  $T$  — in both the clamp bounds and  $n(T, \Delta)$  — that its own constants list never declares; (viii) heavily overloaded symbols ( $\delta, k, n, \lambda$  each carry two or more meanings across sections); (ix) Table 1’s 1,275× improvement entry, which its own SPV and FlyClient sizes contradict ( $496,094/416 \approx 1,193\times$ ; the printed ratio matches a FlyClient size of  $\sim 389$  KB) while the other three entries recompute correctly. The acknowledgements record that CCS reviewers found “Bahack style attacks and problems with the security proof in a previous version” — the variable-difficulty machinery of §3.5 is post-review repair work, which explains some of the visible seams.

## 4 The FlyClient protocol

### 4.1 The model, and what the adversary may do

The analysis lives in the variable-difficulty Bitcoin backbone model of Garay–Kiayias–Leonardos. Execution proceeds in rounds; each active player makes  $q$  sequential random-oracle queries per round; a block  $B$  with target  $T$  is valid iff  $H(B) < T$ ; broadcast messages arrive next round, with the adversary free to reorder deliveries. Honest participation  $n = \{n_r\}$  is assumed  $(\gamma, s)$ -respecting — in any window of fewer than  $s$  rounds, honest power varies by at most a factor  $\gamma$  — and targets are recalculated once per epoch of  $m$  blocks: with the last  $m$  blocks mined in  $\Delta$  rounds, the raw retarget  $T' = T\Delta f/m$  (aiming at  $m$  blocks per  $m/f$  rounds) is clamped so that the target moves by at most a factor  $\tau$  per epoch, the *dampening filter* (ePrint 2019/226, Definition 1, §3.1, p. 7; the paper notes Bitcoin operates with  $\tau = 4, m = 2016, f = 0.03$ ). The clamp is not a tuning detail: it is what makes the difficulty-raising attack of §2.2 survivable, and every feasibility check of §3.5 is an image of it. Backbone parameters are chosen so persistence and liveness hold, so a single chain is held by all honest full nodes.

The adversary is *rushing* (each round it sees all honest messages before choosing its own), mines with at most a  $\mu$  fraction of honest power, and wants the light client to accept a chain of at least the honest chain’s cumulative difficulty. The verifier knows only the genesis block. Two structural facts already constrain the attack. A fork can carry no honest block from after its fork point — the MMR roots inside those blocks would not match the fork (Corollary 4, §3.3) — and no adversarial work from before the fork point can be re-used after it, since the MMR forces each block to be created after everything before it. What remains is the assumption the whole analysis is priced in:

**Assumption 2** ( $(c, L)$ -adversary; ePrint 2019/226, §3.2, p. 8). *There exists no adversary in the variable-difficulty model that respects the target recalculation function from Definition 1 and can, with non-negligible probability, produce a fork that contains more than  $L$  blocks such that a  $c$  fraction of the difficulty weight in these blocks is honest.*

**Remark 4.1** (Reading Assumption 2, and reading  $c$ ). The printed clause “a  $c$  fraction ... is honest” is loose; the meaning used throughout the paper’s own analysis is a cap on the adversary: in any fork of length at least  $L$ , at most a  $c$  fraction of the fork’s difficulty weight can be validly mined, so a fork matching the honest chain’s weight is at least a  $(1 - c)$  fraction fake. Separately,  $c$  is the ratio of adversarial to *honest* mining power — an adversary at  $c = 0.9$  holds  $c/(1 + c) \approx 47.3\%$  of the network’s total (§7.1, p. 20). Short forks are deliberately outside the assumption: for small  $L$  the variance of mining lets any adversary win occasionally, which is why  $L$  trailing blocks are always checked in full. Connecting  $(c, L)$  formally to the backbone parameters  $(\mu, \gamma, s, \dots)$  is conjectured possible and left open by the paper; the constant-difficulty case is the one bridge actually proved:

**Lemma 1** (ePrint 2019/226, p. 8). *In the constant-difficulty backbone, let  $X$  be the number of blocks mined by any adversary while the honest chain adopts  $L$  blocks, the adversary’s rate at most  $\mu$  times the honest adoption rate. Then for  $c > \mu$ ,*

$$\Pr[X \geq cL] \leq e^{L(c-\mu)} \cdot (c/\mu)^{-cL}.$$

The proof is a Chernoff bound on a dominating Poisson variable with mean  $\mu L$ , optimised at  $t = \ln(c/\mu)$ ; Corollary 1 then picks, for every  $\mu$  and  $L = \Theta(\lambda)$ , a  $c < 1$  making the right-hand side  $2^{-\lambda}$  — so at constant difficulty the  $(c, L)$ -assumption is a theorem, not an axiom.

## 4.2 What security means here

The object being proved about is the paper’s SPV predicate: for a chain  $C$  ending in  $B_n$ , the predicate holds iff every block carries correct proof of work following the difficulty-adjustment function and each block’s hash is contained in its direct successor’s header (Definition 2, p. 9). Because persistence guarantees nothing about the last  $k$  blocks, proofs for chains differing only in the last  $k$  blocks count as proofs for the same chain, the highest-difficulty one representative. The verifier is a light client knowing only genesis, with oracle access to  $H$  for checking work; it receives proofs from several provers, *at least one honest* — an assumption that deliberately prices eclipse attacks out of scope — and accepts the predicate for the head of greatest cumulative difficulty. A proof protocol is *secure* if the verifier’s accepted head shares a common prefix (up to the last  $k$  blocks) with the chain of every honest full node (Definition 3), and *succinct* if every proof any efficient prover can emit has size  $O(\text{polylog}(N))$  in the honest chain length  $N$  (Definition 4, inherited from the NIPoPoW paper). FlyClient is thus a NIPoPoW in the primitive sense; “NIPoPoW” as a protocol name usually means the superblock construction of §2.3, which FlyClient replaces.

## 4.3 The interactive game

The interactive protocol (Algorithms 1–2, p. 11) is short enough to give in full. Each prover claims a chain of  $n + 1$  blocks and opens with its head, whose header carries the MMR root over the first  $n$  blocks. If all provers present the same head, the client is done. Otherwise the client samples  $k$  heights per prover from the distribution of §4.4; for each sampled height the prover returns that block’s header with its MMR inclusion path, and the client runs the per-sample check set of §3.4: the path folds to the head’s root; the Get\_Root fold of the same path equals the root stored inside the sampled header; the sampled header’s proof of work is valid, at a target the difficulty-MMR metadata declares and constrains (§3.5). Any failure rejects the prover’s chain; a chain surviving all samples is accepted. Provers claiming different lengths are handled by the NIPoPoW generic verifier, longest claim first.

Two consequences of the commitment structure do quiet work here. First, transaction verification decouples from synchronisation: once the head is accepted, membership of an old block is one MMR inclusion proof against the head’s root, and membership of a transaction is one ordinary SPV Merkle proof inside that block — roughly 1.5 KB on Ethereum parameters against ~15 KB for superblock NIPoPoW (§1, p. 4). Second, the unstable tip is harmless: even a maliciously chosen head cannot commit to invalid blocks or omit stable ones — its MMR must contain every stable block of the honest chain, or sampling exposes it — so the client may run the protocol against the freshest head and still learns the stable prefix correctly; it stores one recent *stable* block between executions (§4.2, pp. 11–12).

## 4.4 From uniform sampling to the optimal distribution

The sampling distribution is the information-theoretic heart of the protocol, and the paper builds it in three failed-then-fixed steps worth keeping, because each failure names a constraint the final distribution satisfies.

**Uniform sampling fails on short forks.** If the fork point  $a$  were known, every sample past it would be invalid with probability  $1 - c$  and  $k$  samples would catch the fork with probability  $1 - c^k$ . But the verifier does not know  $a$ , and a fork near the tip hides from uniform samples: almost all land before the fork, where the adversary’s chain is honestly valid (§5.1, p. 12).

**Binary search is interactive.** With one honest prover guaranteed, the verifier can binary-search the first height at which two provers disagree —  $\log n$  adaptive queries — and then sample  $k$  blocks past the fork point. This works but is multi-round and adaptive, so it cannot be made non-interactive by Fiat–Shamir; FlyClient needs a one-shot, public-coin protocol (§5.2, p. 13).

**Dyadic suffixes get the right shape.** The one-shot fix oversamples the tip: for every  $j \in [0, \log n)$ , sample  $k$  blocks uniformly from the last  $n/2^j$  blocks, and once the interval is at most  $\min(L, k)$  blocks, check all of them.

**Lemma 2** (ePrint 2019/226, p. 13). *With  $k \log n$  samples, the probability that the verifier fails to sample any invalid block is at most  $((1 + c)/2)^k$ .*

The proof needs only one of the  $\log n$  intervals: whatever the fork point, some dyadic interval of size  $n'$  has it in its first half, so more than  $n'/2$  of that interval is post-fork and at least  $(1 - c) \cdot n'/2$  of it is fake;  $k$  uniform samples from that interval alone miss with probability at most  $((1 + c)/2)^k$ . The analysis discards every other interval’s samples — the paper’s own stated question, “can we achieve a better bound?”, is what the optimal distribution answers.

**The optimal distribution.** Reordered, the dyadic protocol is  $q$  i.i.d. draws from a staircase density; the right question is which density maximises the single-draw probability  $p$  of hitting a fake block, against an adversary who places its validly-mined blocks adversarially — then  $q$  draws catch with probability  $1 - (1 - p)^q$ . Model the chain as the continuum  $[0, 1]$ , genesis at 0, head at 1. Two structural facts pin the answer. A non-increasing density is never uniquely optimal — an adversary shifts fakes toward the lower-probability late interval, so mass should increase toward the tip (Lemma 3, pp. 14–15). Against an increasing density, the optimal adversary forking at  $a$  puts all its validly-mined weight at the end of its fork, so the segment  $[a, 1 - (1 - a)c]$  is entirely fake and the single-draw catch probability at fork point  $a$  is  $\int_a^{1+ac-c} f(x) dx$  (normalised). Maximising the minimum over  $a$  forces indifference — the same catch probability at every fork point — and differential analysis yields

$$f(x) = \frac{1 - c}{c(1 - x)}, \quad \int_a^{1+ac-c} f(x) dx = \frac{(c - 1) \ln c}{c} \text{ for every } a,$$

an integral independent of  $a$ , as required. The density diverges at the tip —  $\int_0^1 f = \infty$  — and the repair is the protocol’s signature move: restrict sampling to  $[0, 1 - \delta]$  and *always check the last  $\delta$  fraction of the chain in full*. The normalised sampling density is

$$g(x) = \frac{1}{(x - 1) \ln \delta}, \quad x \in [0, 1 - \delta],$$

(both factors negative, so  $g > 0$ ), and the single-draw catch probability becomes  $p = \log_\delta(c)$ : choosing  $\delta = c^k$  gives  $p = 1/k$  exactly, while any fork point past  $(c - \delta)/c$  pushes the fake segment into the always-checked suffix and is caught with probability 1.

**Theorem 2** (Optimal sampling distribution; ePrint 2019/226, p. 16). *Given that the last  $\delta = c^k$  fraction of the chain contains only valid blocks and the adversary can create at most a  $c$  fraction of valid blocks after the fork point, the sampling distribution with PDF  $g(x) = 1/((x - 1) \ln \delta)$  maximises the probability of catching any adversary that optimises the placement of invalid blocks.*

Optimality is an averaging argument: any candidate density must give catch probability at least  $p^*$  to each of the  $k$  fork points  $a = 1 - c^i$ , and the corresponding intervals tile the support, so  $1 \geq kp^*$  and no density beats  $1/k$ . The sample count follows from  $p_m = (1 - 1/k)^m \leq 2^{-\lambda}$ :

$$m \geq \frac{\lambda}{\log_{1/2}(1 - 1/\log_c(L/n))}, \quad m \approx \lambda \log_c(1/2) \ln(n) = O(\lambda \log_{1/c} n),$$

where the always-checked suffix of  $L = \delta n$  blocks fixes  $k = \log_c(L/n)$ . The suffix length trades against the sample count and is bounded below by the  $(c, L)$ -assumption itself; the implementation optimises it numerically (about a hundred blocks at Ethereum parameters, Figure 4). Each sampled block costs a header plus a  $\log_2(n)$ -hash Merkle path, giving:

**Corollary 2** (ePrint 2019/226, p. 17). *Under the  $(c, L)$ -adversary assumption for any constant  $L$ , and using a collision-resistant hash function, the FlyClient proof size is  $\Theta(\lambda \log(n) \log_{1/c}(n) + L)$ .*

Unlike the superblock protocols, whose succinctness degrades to full SPV under adversarial influence, this bound holds against *all* adversaries.

#### 4.5 Removing the verifier

The protocol is public-coin — the verifier contributes only randomness — and one-shot, so Fiat-Shamir applies: the sampling randomness is derived by hashing the head of the chain with an ordinary secure hash (the paper suggests SHA-3), and a verifier of the transcript checks the derivation along with the samples (§6.2, p. 19). Grinding is priced in proof-of-work: a cheating prover obtains fresh randomness only by re-mining the head, and the assumption bounds its total puzzle budget by  $c \cdot n$ , so a union bound makes the non-interactive protocol secure whenever the interactive soundness satisfies

$$p_m < \frac{2^{-\lambda}}{c \cdot n}.$$

Two deployment properties fall out (§8.1, p. 21). Proofs are *transferable*: anyone can relay a proof without recomputation. And for a given chain the proof is *unique* — the randomness is a function of the head — so one party can compute the proof for the entire network.

#### 4.6 The theorem, and staying synchronised

**Theorem 1** (FlyClient; ePrint 2019/226, pp. 10, 20). *Assuming a variable-difficulty backbone protocol such that all adversaries are limited to be  $(c, L)$ -adversaries as per Assumption 2, and assuming a collision-resistant hash function  $H$ , the FlyClient protocol is a secure NIPoPoW protocol in the random-oracle model as per Definition 3 with all but negligible probability. The protocol is succinct with proof size  $O(L + \lambda \log_{1/c}(n) \log_2(n))$ .*

The proof is the composition of everything above, in five steps: Assumption 2 speaks only of retargeting-respecting adversaries, and Lemma 4 (§3.6) converts every other adversary into one at no loss; Merkle soundness plus the prefix-extension corollaries make the MMR position- and weight-binding, so sampled answers are pinned to one chain; Theorem 2 with Corollary 2’s parameters makes evasion of sampling negligible; Fiat-Shamir is secure for one-round public-coin protocols in the random-oracle model; and the transcript —  $L$  suffix blocks plus  $O(\lambda \log_{1/c} n)$  sampled blocks, each with a  $\log_2(n)$ -hash path — has the stated size.

A synchronised client also stays synchronised cheaply. Theorem 3 (Subchain proofs, §8.2, p. 22): a client that accepted a chain of length  $n$ , later handed a proof for the extension from  $n$  to  $n'$  together with one MMR proof that its block  $n$  lies in block  $n'$ ’s commitment, accepts nothing it would not have accepted from a full  $n'$ -proof — a fork after  $n$  faces a fresh FlyClient instance with block  $n$  as genesis, and a fork before  $n$  fails the prefix proof outright. In practice no custom subproof is needed: the full proof is transferable, so a returning client verifies only the part past its checkpoint, and servers can precompute checkpoint proofs at chosen heights.

#### 4.7 What the evaluation showed

The implementation created and verified every tested proof in under a second. At Ethereum’s 2019 parameters ( $c = 1/2$ ,  $\lambda = 50$ ,  $n$  up to seven million), the numerically-optimised trailing suffix stays near a hundred blocks, the sample count near 400–470, and the whole proof under 530 KB — against

3.4 GB of SPV headers, the 6,627× of Table 2 — with the paper’s “MMR proof optimization” (unexplained in the text, but plausibly deduplication of shared path nodes across samples) accounting for roughly a third of the size (§7, Figures 4, 6). One curiosity in the Ethereum data is instructive: proof size *falls* between three and four million blocks because the difficulty bomb concentrated so much weight near the tip that the always-checked suffix covered a larger weight fraction, cutting the sample count — difficulty growth helps the sampler. Two caveats frame all the numbers.

**Remark 4.2** (Ethereum is outside the theorem). Ethereum’s moving-average difficulty rule does not fit the variable-difficulty backbone model, so Theorem 1 does not cover it; the paper is explicit that running FlyClient on Ethereum is possible “only with heuristic security guarantees” (§3.1, pp. 7–8; §7, p. 20). The point recurs for Zcash in §8: whether Zcash’s own averaging-window retarget (distinct from both Bitcoin’s and Ethereum’s) satisfies the model is exactly the kind of question a Zcash FlyClient ZIP would have to answer, and none exists.

The second caveat is historical and recorded in the paper’s own acknowledgements: CCS reviewers found “Bahack style attacks and problems with the security proof in a previous version” — so the variable-difficulty analysis is repair work, added after review. Finally, the paper closes the circle to proofs of sequential work: the chain construction is essentially Cohen–Pietrzak PoSW with block hashes as leaves, and their verifier is FlyClient with a uniform distribution — a FlyClient chain *is* a PoSW, with the stronger guarantee that no point of the chain has a constant fraction of corrupted leaves (§8.4, pp. 23–24).

## 5 ZIP-221: the deployed chain-history tree

Zcash deployed the server half of FlyClient as consensus. ZIP-221 — title “FlyClient — Consensus-Layer Changes”, Status Final, deployed with the Heartwood upgrade — replaces the block header’s Sapling treestate commitment with a commitment to the root of a Merkle mountain range over the chain’s history, each node of which carries the metadata a difficulty-sampling verifier needs. The header field’s lineage tells the story compactly: `hashReserved` before Sapling, `hashFinalSaplingRoot` from Sapling, renamed `hashLightClientRoot` by ZIP-221 at Heartwood, renamed again `hashBlockCommitments` by ZIP-244 at NU5 (§6). The stated motivation is the paper’s protocol verbatim: MMR proofs are to let light clients verify “the validity of a block chain received from a full node”, block inclusion, and “certain metadata of any block or range of blocks in that chain” (ZIP 221, §“Motivation”), with header downloads dropping from linear to logarithmic.

### 5.1 Epoch scoping: one tree per network upgrade

The single largest structural departure from the paper is what the tree covers. The paper’s MMR runs from genesis; ZIP-221’s does not:

“The protocol requires that an MMR that commits to the inclusion of all blocks since the preceding network upgrade ( $B_x, \dots, B_{n-1}$ ) is formed for each block  $B_n$ . The root  $M_n$  of the MMR MUST be included in the header of  $B_n$ .” (ZIP 221, §“Motivation”;  $x$  is the activation height of the preceding network upgrade.)

Three consequences follow from the exact definition (§“Tree nodes and hashing”). First, the commitment is delayed one block: the root inside  $B_n$  covers leaves through  $B_{n-1}$  only, so a block never commits to itself — the ZIP’s rationale notes the corollary that a block’s final Sapling root now surfaces one block later than it did pre-Heartwood, and accepts it because anchoring on the most recent treestate “is already unsafe in reorgs”. Second, the tree *resets at every network upgrade*: the preceding-upgrade height  $x$  is the last activation height strictly below  $n$ , so the block at an activation height  $A$  carries the completed previous epoch’s root, and the very next block  $A + 1$  already commits to the new epoch’s single-leaf tree over  $B_A$ . Every consensus branch thus owns a self-contained tree

— which composes with the personalisation scheme below into cross-branch domain separation — and any future Zcash FlyClient must chain per-epoch proofs across upgrade boundaries, a composition the ZIP does not discuss. Third, the genesis and activation edge cases are handled by sentinel: `hashChainHistoryRoot` is undefined at genesis, and in the block that activates the ZIP the header field is all zero bytes, which “MUST NOT be interpreted as a root hash”.

## 5.2 Numbering, peaks, and bagging

Where the paper specifies its MMR as a single-rooted recursion (Definition 3.1) and never says “peak”, ZIP-221 specifies the same object in the classical presentation the paper omitted (Remark 3.2), adapted from Grin’s: nodes are numbered in *creation order* — each leaf followed immediately by every internal node its arrival completes — and the forest of perfect subtrees (“mountains”, leftmost peak always the highest) is *bagged* into a single root by repeatedly joining the two *leftmost* peaks: a left fold, so peaks  $P_1, P_2, P_3$  bag as  $(P_1 || P_2) || P_3$ . Navigation is arithmetic on positions: a node at position  $p$  with altitude  $h$  has its right sibling at  $p + 2^{h+1} - 1$  and its left child at  $p - 2^h$ , with the ZIP’s own caveat that bagging nodes fall outside these rules. The ZIP’s eleven-leaf worked example — 19 nodes at positions 0–18, peaks at positions 14, 17, 18 with altitudes 3, 1, 0 — and both navigation identities reproduce exactly under this volume’s independent implementation (script-verified 2026-07-19). Reorgs run the construction backwards: deleting the rightmost leaf pops the right spine and re-bags,  $O(\log k)$  in the right subtree’s leaf count.

## 5.3 Node metadata: the complete field list

Every node — leaf, internal, root, treated uniformly — carries the same field vector, serialised in the exact order of Table 3. A leaf takes both members of each earliest/latest pair from its own block; an internal node inherits *earliest*-fields from its left child, *latest*-fields from its right child, and sums the additive fields. The ZIP splits the vector by purpose: seven fields — the timestamp, target, and height pairs, plus the cumulative work — are the “FlyClient Requirements”, chosen “via discussion with an author of the FlyClient paper” to let a light client reason about the difficulty-adjustment algorithm over a block range, exactly the role Definition 7 of the paper assigns its  $(w, t, D_{\text{start}}, D_{\text{next}}, n)$  payload (§3.5); the remaining fields are Zcash-specific light client aids with no counterpart in the paper. Among the latter, the shielded transaction counts are the interesting design: they let a client reliably learn whether a malicious server is withholding shielded transactions, and let it skip header ranges containing none — the ZIP notes the dominant light-client cost is transmitting Equihash solutions in headers. An earlier draft counted Sprout and Sapling together; it was rejected because a Sapling-only client could not distinguish a withheld Sapling transaction from an unrequested Sprout one, and the ZIP fixes the pattern that every future pool gets its own count — which NU5’s Orchard trio and the draft Ironwood trio then follow (§6).

Serialised size follows from the widths: the V1 fixed portion is  $32 + 4 + 4 + 4 + 4 + 32 + 32 + 32 = 144$  bytes plus three CompactSize integers of 1–9 bytes each, giving the ZIP’s stated 147–171 bytes; NU5 adds 64 fixed bytes and one more CompactSize, giving 212–244 (script-verified against both stated ranges). Two of the ZIP’s own field notes are worth carrying: `nLatestTimestamp` may be *smaller* than `nEarliestTimestamp` in subtrees narrower than `PoWMedianBlockSpan`, because header timestamps are not consensus-monotonic at that scale; and within one branch, the leaf hash of  $B_{n-1}$  equals `hashPrevBlock` of the following header exactly, so the tree adds no new hashing at the leaves.

## 5.4 Hashing and personalisation

All internal hashing is BLAKE2b-256 with the 16-byte personalisation `ZcashHistory || CONSENSUS_BRANCH_ID` — the ASCII prefix’s twelve bytes followed by the branch ID’s four little-endian bytes, so for Heartwood’s branch `0xF5B9230B` the trailing bytes are `0B 23 B9 F5`. Baking the branch ID into the personalisation means “valid nodes from one consensus branch cannot be used to make

| #  | Field                   | Leaf value                                                                   | Internal rule                                                      |
|----|-------------------------|------------------------------------------------------------------------------|--------------------------------------------------------------------|
| 1  | hashSubtreeCommitment   | the block’s consensus block hash, internal byte order, <i>unpersonalised</i> | personalised BLAKE2b-256 of the two children’s full serialisations |
| 2  | nEarliestTimestamp      | header nTime                                                                 | inherit left                                                       |
| 3  | nLatestTimestamp        | header nTime                                                                 | inherit right                                                      |
| 4  | nEarliestTargetBits     | header nBits                                                                 | inherit left                                                       |
| 5  | nLatestTargetBits       | header nBits                                                                 | inherit right                                                      |
| 6  | hashEarliestSaplingRoot | final Sapling treestate root                                                 | inherit left                                                       |
| 7  | hashLatestSaplingRoot   | final Sapling treestate root                                                 | inherit right                                                      |
| 8  | nSubTreeTotalWork       | $\lfloor 2^{256} / (\text{ToTarget}(\text{nBits}) + 1) \rfloor$              | sum (modulo $2^{256}$ )                                            |
| 9  | nEarliestHeight         | block height                                                                 | inherit left                                                       |
| 10 | nLatestHeight           | block height                                                                 | inherit right                                                      |
| 11 | nSaplingTxCount         | count of transactions with non-empty vSpendsSapling or vOutputsSapling       | sum                                                                |
| 12 | hashEarliestOrchardRoot | final Orchard treestate root                                                 | inherit left                                                       |
| 13 | hashLatestOrchardRoot   | final Orchard treestate root                                                 | inherit right                                                      |
| 14 | nOrchardTxCount         | count of transactions with non-empty vActionsOrchard                         | sum                                                                |

Table 3: The ZIP-221 node field vector, in serialisation order. Fields 1–11 are the V1 vector (Heartwood onward); fields 12–14 join at NU5 and are omitted before it. Hash fields are `char[32]`, timestamps and target bits `uint32`, work `uint256`, heights and counts `CompactSize` integers. The seven FlyClient-requirement fields are 2–5 and 8–10. Sources: ZIP 221, §“Tree Node specification”; ZIP 252.

false statements about parallel consensus branches” (ZIP 221, §“Rationale”) — composed with epoch scoping, each upgrade’s tree is domain-separated from every other’s. Two boundaries of the scheme are exact: a *leaf*’s hash field is the block’s consensus block hash itself, unpersonalised; and an *internal* node hashes the concatenation of its children’s full canonical serialisations — metadata and all, not merely their digests — so every aggregate field is hash-bound at every level. The header commitment is one more hash of the same kind: `hashChainHistoryRoot` is the personalised BLAKE2b-256 digest of the *serialised root node*, so the 32-byte header field commits to the whole epoch’s aggregates — total work, boundary timestamps and targets, boundary treestate roots, shielded transaction counts — not only to the leaf set.

**Remark 5.1** (Which branch ID, at an activation block). The ZIP’s prose sets the personalisation branch ID to that of “the epoch of the block containing the commitment”, but its pseudocode hashes with the branch ID carried by the *nodes* (`make_parent` uses the left child’s, `make_root_commitment` the root’s). Within one epoch-scoped tree the two readings coincide everywhere except at an activation block, whose header carries the *previous* epoch’s finished root: there the pseudocode — and the ZIP’s own summary that “each node in the tree commits to the consensus branch that produced it” — resolve the tension in favour of the tree’s own branch ID, and the deployed implementations agree (§7).

## 5.5 The consensus rules at Heartwood

The consensus delta is three rules on one renamed field (ZIP 221, §“Block header semantics and consensus rules”): before activation, the Sapling rule persists under the new name (`hashLightClientRoot` is the final Sapling treestate root); in the activation block, the field **MUST** be all zero bytes, not interpreted as a root; in every subsequent block, the field **MUST** equal `hashChainHistoryRoot`. The header’s byte format and version are untouched — a deliberate boundary, argued in the rationale: ZIP-301 obliges miners to ignore jobs with unknown header versions, so a version bump risks mining-hardware breakage for a change that is semantically invisible to miners. The rationale also disposes of two alternatives: overloading is safe against a



miner contriving a Sapling treestate whose root collides with an MMR root (a preimage over 32 layers of Pedersen hashing), and the seemingly cleaner home for the root — replacing `hashPrevBlock`, exactly the paper’s suggestion (§3.4) — was rejected because it would require “massive refactoring” and hurt performance, reorgs in particular being “fragile, performance-critical” and reliant on backwards iteration over the chain history.

## 5.6 Deltas from the paper’s construction

For a reader holding both objects, the differences reduce to six that matter and a summary judgement. The tree is epoch-scoped, not genesis-rooted, with a one-block commitment delay (§5.1). Node metadata is richer than Definition 7 of the paper: the seven `FlyClient` fields realise the paper’s  $(w, t, D_{\text{start}}, D_{\text{next}}, n)$  payload — with both ends of each pair carried explicitly rather than one derived — and the Sapling/Orchard fields have no paper counterpart. Hashing is personalised, branch-separated BLAKE2b-256 over full child serialisations, where the paper hashes bare child values with an unspecified generic  $H$ . The root reaches the header through an existing repurposed field rather than a new one — and, from NU5, only as the first input of an outer commitment (§6). The bagging order is fixed left-fold-leftmost (§5.2) — the opposite of the right fold the paper’s recursion implies (Remark 3.2), so the two constructions’ bagged roots differ whenever the tree has three or more peaks. The ZIP’s own summary is accurate: “Other than the metadata commitments, the MMR tree’s construction is standard”, and its append algorithm is “adapted from” the paper’s.

**Remark 5.2** (Defect ledger for a Final ZIP). The ZIP’s illustrative pseudocode — explicitly non-normative, but the only executable statement of intent the ZIP contains — has at least six defects in the checked-out text: the node class declares `nSaplingTxCount` twice where the second must read `nOrchardTxCount`; `serialize` and `get_peaks` reference instance fields without `self`; `node`. `qualification`; `make_parent` reads `sapling_root`/`orchard_root` attributes that exist on the block type, not the node type; `append` and `delete` read `latest_height`/`earliest_height` where the node fields are named `nLatestHeight`/`nEarliestHeight`; and a stray Rust-ism (`peaks.push`) survives in the Python. None affects the normative field table or consensus rules; all are trivially repaired by the reference implementations (§7). Separately, no ZIP-221 test vectors exist anywhere in the zips repository.

**Remark 5.3** (The ZIP’s own security caveat). ZIP-221 is unusually candid about the limits of its security story, and the caveat is load-bearing for §8: the `FlyClient` paper “only analyses these security properties in chain models with slowly adjusting difficulty, such as Bitcoin”, leaves rapidly-adjusting chains — “such as Zcash or Ethereum” — as an open problem with “only heuristic security guarantees”, and although the difficulty-reasoning fields were chosen with a paper author, “the use of these fields has not been analysed in the academic security literature”. The ZIP concludes that “a more in-depth security analysis of `FlyClient` should be performed before designing a `FlyClient`-based light client ecosystem for Zcash”. This is the same boundary Remark 4.2 drew from the paper’s side: Zcash’s averaging-window retarget stands outside the proven model, and nobody has closed the gap in either direction since.

## 6 NU5: the commitment becomes one of two

### 6.1 Why an authorising-data root exists

NU5’s ZIP-244 made transaction identifiers non-malleable by excluding witness data — proofs and signatures — from the txid digest. The repair creates a gap: `hashMerkleRoot`, now built over non-malleable txids, no longer commits to the whole serialised transaction. ZIP-244 fills it with a parallel commitment: each transaction’s *authorising data commitment* digests exactly the excluded material (transparent input scripts; Sapling proofs and signatures; Orchard proofs, spend-authorisation signatures, and binding signature) under fixed BLAKE2b-256 personalisations, gathered by an

outer per-branch digest, with pre-v5 transactions assigned the placeholder of 32 0xFF bytes; and `hashAuthDataRoot` is the root of a padded binary Merkle tree (personalisation `ZcashAuthDatHash`, null leaves `[0u8; 32]`) over those commitments, in insertion order identical to the txid tree, so one path position names the same transaction in both trees. The pair (txid tree, auth tree) restores full commitment to every byte a block carries.

## 6.2 hashBlockCommitments

Rather than widen the header, ZIP-244 overloads the ZIP-221 field a second time. From NU5 the field is named `hashBlockCommitments` and holds the personalised digest

$$\text{BLAKE2b-256}_{\text{ZcashBlockCommit}}(\text{hashLightClientRoot} \parallel \text{hashAuthDataRoot} \parallel [0\text{u8}; 32]),$$

whose first input is the ZIP-221 chain-history root, second the authorising-data root, and third a 32-zero-byte terminator. The chain-history root therefore no longer appears directly in any NU5 header; a client checking it must carry the auth-data root (or its subtree) alongside any MMR proof. The terminator is deliberate architecture: the ZIP describes the value as “a linked list of hashes terminated by arbitrary data”, so a future upgrade can replace the zeros with the hash of a further structure ending in its own terminator; full validators must always use the entire structure of the latest *activated* protocol version they support, and servers for older light clients SHOULD supply the replacement value so old verification equations keep closing. Two activation details differ from Heartwood’s: there is *no* zero-byte sentinel at NU5 — the activation block already uses the new derivation, its inner chain-history input being the completed Canopy-epoch root — and the same upgrade extends the node vector with the Orchard trio (Table 3, fields 12–14), following the one-count-per-pool pattern. The draft ZIP-258 continues both patterns for NU6.3: an Ironwood trio (fields 15–17, aggregated identically to the Orchard trio, node size 277–317 bytes — the arithmetic checks against the field widths), with `hashChainHistoryRoot` changing “only through the added node fields” and `hashBlockCommitments` otherwise unaffected. That trio is the one this series’ Ironwood volume already cites at its turnstile (*Ironwood Guide*, §“The turnstile and migration”). ZIP-258 remains Draft as a document, but these rules and fields are active consensus and are implemented by Zebra; the ZIP proposes that `zcashd` will not implement NU6.3.

## 6.3 Activation constants

| Upgrade   | Branch ID  | Mainnet   | Testnet   | Node vector          |
|-----------|------------|-----------|-----------|----------------------|
| Heartwood | 0xF5B9230B | 903,000   | 903,800   | V1 (fields 1–11)     |
| NU5       | 0xC2D6D0B4 | 1,687,104 | 1,842,420 | V2 (+ Orchard trio)  |
| NU6.3     | 0x37A5165B | 3,428,143 | 4,134,000 | V3 (+ Ironwood trio) |

Table 4: Chain-history activation constants. Sources: ZIP 250, ZIP 252, ZIP 258 (Draft). The NU5 testnet height is the second activation: a first testnet NU5 at height 1,599,200 under branch 0x37519621 (an earlier Orchard circuit) was rolled back, the testnet forking from height 1,599,199.

# 7 Deployed implementations

Three codebases realise ZIP-221: the `zcash_history` crate (“to be used in zebra and via FFI bindings in `zcashd`”, its own words), Zebra’s wrapper and state service, and `zcashd`’s C++ cache with a Rust FFI bridge. Zebra pins version 0.5.0 of the crate; `zcashd` pins 0.4.0 — a skew with consequences registered in §7.4.

## 7.1 The `zcash_history` crate

The crate stores the MMR in its *array representation*: leaves and completed internal nodes in append order, addressed by u32 index. Two link kinds separate what persists from what does not (`src/lib.rs:42--49`): `Stored(u32)` indexes the persistent array; `Generated(u32)` indexes an ephemeral per-instance vector holding the *bagging* nodes that join unequal peaks into a root — and serialisation refuses entries with `Generated` children (`src/entry.rs:98--100`), so exactly ZIP-221’s node array persists, never the bag. An entry serialises as a one-byte tag (node or leaf), two little-endian u32 child indices for nodes, then the node data (`src/entry.rs:69--106`); the node data is ZIP-221’s field vector in ZIP order, `CompactSize` integers included, with maxima asserted in tests at exactly the ZIP’s figures — 171 bytes (V1), 244 (V2), 317 (V3, Ironwood), entry maximum  $317 + 9 = 326$  (`src/node_data.rs:463--482`). One deliberate absence: the consensus branch ID is *context, not data* — it is never serialised, and every deserialiser must supply it (`src/node_data.rs:135--139`).

Hashing matches §5.4 exactly (`src/version.rs:21--26`, `43--65`, `102--106`): the personalisation is `ZcashHistory || LE32(branch ID)`; a parent’s commitment hashes the concatenation of both children’s *full serialisations*; and the hash of the root node’s serialisation is `hashChainHistoryRoot`. A dedicated test (`v3_combine_hash_commits_to_ironwood_fields`) pins the property that parent hashes commit to child metadata, and a conformance module walks all three node versions against the `zcash-test-vectors` ZIP-221 vectors, generated by an independent Python reimplementation (`src/test_vectors.rs:1--13`).

The `Tree` type is explicitly a *view*, not a container (`src/tree.rs:5--15`): it is instantiated per operation from a caller-supplied peak set — `Tree::new(length, peaks, extra)`, which inserts the peaks and left-folds them into generated bagging nodes — used for a few appends or truncations, and dropped. Appending (`src/tree.rs:137--194`) pushes the leaf, collects peaks by descending from the root and stopping at complete subtrees, merges equal-sized peaks right-to-left into *stored* nodes, and re-bags the remainder left-to-right into *generated* nodes; truncation runs the inverse, consuming the caller-supplied extra nodes down the right spine. Completeness itself is derived from metadata, not shape bookkeeping: a subtree is complete iff `end_height - (start_height - 1)` is a power of two (`src/entry.rs:38--45`) — an elegant economy with a sharp edge registered below. Notably, the crate never computes which array positions are peaks; every caller reimplements the walk — the crate’s own example 1-based, Zebra and `zcashd` in a matching 0-based form (highest altitude  $\lfloor \log_2(n+1) \rfloor - 1$ , first candidate  $2^{\text{alt}+1} - 2$ , descend left while past the end, then hop right siblings) — and the same fourteen-leaf, twenty-five-node ASCII diagram annotates the walk in both Zebra and `zcashd`, Zebra’s comment attributing the code to the crate example and the explanation to `zcashd`.

## 7.2 Zebra

Zebra’s wrapper (`zebra-chain, src/primitives/zcash_history.rs`) builds leaves from blocks: branch ID from the block’s network upgrade, subtree commitment the block hash, both ends of each pair the block’s own time, target, treestate roots and height, work from the block’s difficulty threshold, and the pool transaction counts — with a named `BlockCommitmentTreeRoots` struct whose stated purpose is that the Orchard and Ironwood roots, which share a Rust type, “cannot be silently swapped, which would corrupt the ZIP-221 chain-history commitment” (lines 23–34). Above it, `NonEmptyHistoryTree` (`src/history_tree.rs`) dispatches the node version by upgrade — Heartwood/Canopy to V1, NU5 through NU6.2 to V2, NU6.3/NU7 to V3, panic before Heartwood — and implements the epoch reset with one line: on a push whose network upgrade differs from the tree’s, “This is the activation block of a network upgrade. Create a new tree.” — the value is wholesale replaced by a fresh single-leaf tree (lines 238–247). Every push then prunes: the peak-set walk retains only peak entries, and the in-memory crate tree is rebuilt from peaks alone, neutralising the view type’s unbounded growth. The whole object is wrapped once more as `HistoryTree(Option<NonEmptyHistoryTree>)`, empty before Heartwood — the crate itself cannot represent an empty tree (its constructor panics on an empty peak set), so emptiness lives a

level up. A guard worth quoting sits on push: heights must be consecutive because “librustzcash assumes the heights are correct and corrupts the tree if they are wrong” (lines 224–236) — the completeness-from-heights economy, fenced.

Validation splits across two layers. Structurally, `Commitment::from_bytes` (zebra-chain, `src/block/commitment.rs:108--151`) parses the header field by epoch into the five-armed enum — pre-Sapling reserved, final Sapling root, the all-zero `ChainHistoryActivationReserved` at the Heartwood activation height, `ChainHistoryRoot` through Canopy, `ChainHistoryBlockTxAuthCommitment` from NU5 — rejecting a non-zero activation sentinel outright. Contextually, the state service (zebra-state, `src/service/check.rs:134--222`) compares: for Heartwood/Canopy blocks, the parent chain’s tree hash against the header root; for NU5 onward, it recomputes

```
BLAKE2b-256ZcashBlockCommit(history root || auth_data_root || [0u8; 32])
```

from the parent tree and the block’s own authorising-data tree and compares that. The check runs in the non-finalized state (in parallel with anchor checks and the chain push) and again for checkpointed blocks in the finalized state; zebra-consensus itself never touches the field, a division of labour the code documents — the checkpoint verifier “doesn’t bind the transaction authorizing data”, the state services do. Persistence is minimal by design: one RocksDB column family holding exactly one record — the current tip tree’s peaks — under the unit key `()`, atomically replaced per block write; a previous epoch’s tree exists nowhere in Zebra once its activation is finalized. Reorgs never truncate: the non-finalized state keeps an Arc-shared `HistoryTree` snapshot per height per chain fork, so rolling back is dropping snapshots, and the crate’s truncation machinery goes unused. The mining path closes the loop: `getblocktemplate` responses carry `chain_history_root`, “The FlyClient chain history root as of the end of the chain tip block” (`src/response.rs:530--536`) — producers, too, consume the tree.

### 7.3 zcashd

The `zcashd` node keeps the tree behind an FFI whose Rust side is thin (`src/rust/src/history.rs`): fixed-width byte arrays (244-byte nodes, 253-byte entries — version 0.4.0’s V2 maxima), three entry points (`append`, `remove`, `hash_node`), and a branch-ID dispatch that names Sprout, Overwinter, Sapling, Heartwood, and Canopy for V1 — omitting Blossom, which falls through the *everything else* arm to V2 (harmless: history trees exist only from Heartwood). The C++ side (`src/coins.cpp:384--630`) holds one `HistoryCache` *per epoch*, keyed by consensus branch ID; `PushHistoryNode` preloads the peak set with the same walk as everywhere else, appends through the FFI (at most 32 new nodes per append — the leaf plus at most 31 ancestors, the  $2^{32}$  tree bound made concrete), and `PopHistoryNode` handles reorgs case by case, including the pleasing impossibility proof in an exception: “a history tree cannot have two nodes” — one leaf is length 1, two leaves are length 3. Validation happens in `ConnectBlock` (`src/main.cpp:3505--3614`): the chain-history root is read from the *previous block*’s epoch’s cache — which at an activation block is the previous epoch’s final root, exactly ZIP-221’s rule — and the header field is checked against the Heartwood sentinel, the bare root (pre-NU5), or `DeriveBlockCommitmentsHash`, byte-identical to Zebra’s recomputation. Persistence is the opposite pole from Zebra’s: LevelDB stores every epoch’s *full node array* indefinitely — per-node records keyed by (epoch, index), plus a length and root per epoch — because `zcashd`’s rollback model reopens old epochs’ trees, and its RPC exposes `chainhistoryroot` on `getblock`. A width workaround mirrors Zebra’s: V1 nodes written by pre-NU5 clients occupy 171-byte records that NU5-aware clients read and zero-pad into 244-byte buffers.

### 7.4 Divergence register

The two validators agreed on every consensus-relevant byte for as long as both ran; the register below records where they differ in structure, and one place they would have differed in consensus had both crossed the NU6.3 boundary unchanged.

- *Version skew at NU6.3.* `zcashd`'s 0.4.0 dispatch sends every post-Canopy branch ID to V2 nodes, while Zebra selects V3 (Ironwood fields) for NU6.3 — as the code stands, the two would compute divergent `hashChainHistoryRoot` values from the activation block onward. ZIP-258 resolves the tension by fiat rather than by code: `zcashd` “will not implement” NU6.3, and Zebra is expected to be the sole full validator (§6); the skew is a documented retirement, not a live consensus bug — but the code alone does not say so.
- *Epoch retention.* Zebra destroys a superseded epoch's tree (tip-only, single-record persistence); `zcashd` retained every epoch's full array, because its rollback path could reopen them. Two correct answers to the same question, shaped by each node's reorg model.
- *Truncation asymmetry.* The crate's `truncate_leaf` and extra-node machinery are exercised only by `zcashd`; Zebra's per-height snapshots make deletion a non-event. A bug in truncation would be a `zcashd`-only bug.
- *A skippable check.* `zcashd` verifies the NU5 recomputation only under the `fCheckAuthDataRoot` flag; the pre-NU5 root equality has no such flag. Zebra checks unconditionally on both paths.
- *Fragile economy, fenced differently.* Subtree completeness derived from height metadata means wrong heights corrupt tree shape silently; Zebra fences with a consecutive-height assertion citing exactly this failure, `zcashd` relies on `ConnectBlock` ordering.
- *Edge-case fallback.* Zebra's NU5 check substitutes the 32-zero-byte sentinel as the “previous root” if the NU5-style commitment appears at the Heartwood activation height itself — reachable only on Regtest or configured testnets; `zcashd` has no equivalent arm.

## 8 The deployment gap

### 8.1 Who computes the root, who verifies it

Every Zcash block since Heartwood carries a correct chain-history commitment, because two classes of software maintain the tree. Consensus validators compute and check it: Zebra in the state service and, until its retirement, `zcashd` in `ConnectBlock` (§7). Miners publish it without computing it: both nodes' `getblocktemplate` hand the miner a `defaultroots` object containing `chainhistoryroot` (and, from NU5, `authdataroot` and `blockcommitmentshash`), so template consumers hash what the node computed — the state-service response type behind the template documents the field as “The FlyClient chain history root as of the end of the chain tip block”, and Zebra's built-in miner rides the same RPC. Even this producer-facing surface has had consumer-visible bugs: a `zcashd` comment memorialises that v4.6.0's legacy template fields “were accidentally set to always be the NU5 value” (`src/miner.cpp:462--466`).

Verification, meanwhile, is purely *reconstructive*: a validator maintains the tree itself and compares 32-byte roots. Nobody, in any deployed software, verifies an MMR *proof* — the operation the structure exists to serve. The RPC surface tells the same story from the read side: `zcashd`'s `getblock` exposes `chainhistoryroot` per block, while Zebra — the maintained node — never emits the root for an existing block at all; its `getblock` code carries the literal comment “`chainhistoryroot` would be here. Undocumented. TODO: decide if we want to support it” (`zebra-rpc, src/methods.rs:3971`), and no RPC in either node serves MMR nodes or assembled proofs.

### 8.2 The light-client stack consumes none of it

The deployed light-client protocol was specified before ZIP-221 and never updated toward it: ZIP-307 does not reference ZIP-221 once, and its security model is the one the *Sync Guide* documents at length — the node/proxy combination “is trusted to provide a correct view of the current best chain state”. The negative results are uniform across the stack (all greps 2026-07-19): `lightwalletd`

contains no FlyClient, chain-history, or MMR code and serves no related endpoint; Zaino ships the same gRPC protocol and touches the field only to pass it through — serialising the header’s hashBlockCommitments bytes, and carrying a comment on its header type, “Optional fields may be added for: hashLightClientRoot (FlyClient proofs)...”, a capability anticipated, not implemented; neither zcash\_client\_backend nor zcash\_client\_sqlite nor zcash\_primitives depends on zcash\_history; the Android SDK has no trace. What the wallet stack does instead is exactly the *Sync Guide*’s subject: hash-chain continuity (prevHash and height) and nothing more — ZIP-307’s own text concedes its header-validation section “has been only partially implemented: currently only prevHash is checked”, and a grep of the wallet backend for Equihash or difficulty code returns nothing.

The consequence deserves plain statement: a Zcash FlyClient would not be “SPV, but smaller”. Deployed Zcash wallets perform no proof-of-work verification whatsoever, so a FlyClient verifier would be the *first* PoW verification the wallet stack has ever done — and the seven metadata fields ZIP-221 carries for it, plus the shielded-transaction counts added specifically so light clients could detect withholding (§5.3), have waited six years without a single consumer.

### 8.3 What an actual Zcash FlyClient still needs

- *Prover data availability.* A proof server needs random access to the full node array. Zebra persists peaks only (§7.2); zcashd persists everything a prover needs (§7.3) but exposes none of it and is being retired.
- *Serving endpoints.* No RPC or gRPC anywhere serves an MMR node, an inclusion path, or an assembled proof — the entire wire surface remains to be designed and specified.
- *The sampling client.* No implementation of the verifier loop — weight-proportional sampling under  $g$ , per-sample checks, difficulty-transition feasibility over ZIP-221’s metadata — exists in any wallet or SDK.
- *Specification work.* Open in full: a proof wire format; the Fiat–Shamir transcript definition for Zcash’s header; the composition rule for chaining per-epoch trees across upgrade boundaries (§5.1 — the ZIP is silent); the security analysis Zcash’s own rapidly-adjusting difficulty rule requires (Remark 5.3); and the marriage of FlyClient header verification with ZIP-307 scanning, which replaces the header-chain trust but not the compact-block detection channel.
- *Practicality data.* The one end-to-end datapoint is academic and recent (web-sourced: Perazzo–Capecci, “Catching the Fly”, arXiv 2604.26736, April 2026): the first working Zcash FlyClient prover, built on a zebra fork with four added RPCs and a ~10-hour database backfill to recover the node array Zebra discards; proofs of ~1.2 MiB at three million blocks (~320 KiB with their distillation), with Equihash’s 1,344-byte solutions dominating per-sampled-header cost. The numbers say mobile plausibility and on-chain-verifier marginality — and that the retrofit cost of Zebra’s peaks-only economy is real but bounded.

**Remark 8.1** (An inversion of readiness). The deployment gap has a neat summary: the node that has the prover’s data (zcashd, full per-epoch node arrays in LevelDB) has no future, and the node that has the future (Zebra) has no data. Neither has an interface. Meanwhile consensus keeps paying the feature’s maintenance cost — the NU6.3 upgrade extends the node vector a second time (§6), Zebra already implements it, and still nothing reads the tree.

### 8.4 Elsewhere: velvet paths and inspired bridges

Context from beyond Zcash, all web-sourced. ECC’s Heartwood announcements framed ZIP-221 as “Flyclient support” enabling interoperability — capability language, with no client deliverable then or since; a 2021 community-forum thread asking about FlyClient implementation drew no answers. The follow-up literature validated Zcash’s deployment choice from an unexpected angle: superlight

clients deployed by *velvet* fork — the uncontentious path the paper itself proposed (ePrint 2019/226, §8.3, p. 23) — admit *chain-sewing* attacks, in which portions of independent forks are stitched into a proof (Kiayias–Polydouri–Zindros, AFT 2021); Zcash’s hard-fork deployment forecloses the class, a point the 2026 prover paper makes explicitly. No published attack on FlyClient’s sampling under its stated model was found; the critical literature attacks velvet deployments and the realism of the adversary model, not the distribution. The nearest production relative is FlyClient-*inspired* rather than FlyClient: Harmony’s Horizon bridge design checkpoints MMR roots for an on-chain client — the roots are final by supermajority vote of a fixed validator committee, a Byzantine-fault-tolerant (BFT) protocol safe while fewer than a third of the committee misbehaves arbitrarily, not by probabilistic PoW sampling. Claims that Grin or Beam “use FlyClient” conflate carrying MMR roots in headers (which Mumblewimble chains do natively — the paper’s own observation) with running a sampling verifier, for which no evidence surfaced.

## 8.5 Relation to the frontier

Two frontier designs of this series touch the tree, and per the sibling-ownership convention this volume owns both pointers. The *Tachyon Guide* records a draft amendment to ZIP-221 among Tachyon’s peripheral ZIPs — unmerged, “still under discussion and nothing has been decided yet” in the designers’ words — so the aggregation layer of a Tachyon-era chain may yet become another node field, exactly as Orchard and Ironwood did. The *Crosslink Guide* changes the meaning of the quantity FlyClient proves: a FlyClient transcript establishes cumulative *proof-of-work weight*, while Crosslink’s trailing finality layer makes a prefix of the chain final by BFT assertion — under such a hybrid, a weight-sampling proof of the unfinalized suffix and a certificate for the finalized prefix are different objects with different trust bases, and no analysis of their composition exists in either direction. Neither interaction is more than noted here; both are open in exactly the sense of §8.3’s specification bullet.

## 8.6 Closing the boundary

The volume’s three-bin summary, restated with everything above behind it. *Specified and deployed*: the chain-history tree — epoch scoped, metadata-rich, branch-separated — maintained by every validator and committed in every header for six years, with conformance test vectors and two independent implementations that agree byte for byte. *Designed but unspecified for Zcash*: the FlyClient protocol itself, proved secure in a model that Zcash’s difficulty rule is not known to satisfy — the ZIP says so, the paper says so, and nothing published since closes the question. *Open*: the bridge — data availability, wire format, epoch composition, the sampling client, and the security analysis — of which only an academic prototype exists, on a fork, six years after the consensus layer shipped its half. The chain has kept its side of an agreement whose other party has yet to arrive.